

Project-3 Report Document

Kiran Gowda – 018761559

Karthik Ganduri – 018779902

➔ `data = pd.read_csv("wine_quality.csv",header=None, names = table_header)`
➔ `print(data)` *# Displaying the dataset from the dataset file wine_quality.csv*

	fixedacid	volacid	citricacid	residualsugar	chlorides	freesulfur	\
0	7.0	0.270	0.36	20.70	0.045	45.0	
1	7.2	0.230	0.32	8.50	0.058	47.0	
2	7.2	0.230	0.32	8.50	0.058	47.0	
3	7.0	0.270	0.36	20.70	0.045	45.0	
4	6.3	0.300	0.34	1.60	0.049	14.0	
...	
4896	4.9	0.235	0.27	11.75	0.030	34.0	
4897	6.1	0.340	0.29	2.20	0.036	25.0	
4898	5.7	0.210	0.32	0.90	0.038	38.0	
4899	6.5	0.230	0.38	1.30	0.032	29.0	
4900	6.5	0.240	0.19	1.20	0.041	30.0	

	totalsulfur	density	pH	sulphates	alcohol	quality
0	170.0	1.00100	3.00	0.45	8.8	6.0
1	186.0	0.99560	3.19	0.40	9.9	6.0
2	186.0	0.99560	3.19	0.40	9.9	6.0
3	170.0	1.00100	3.00	0.45	8.8	6.0
4	132.0	0.99400	3.30	0.49	9.5	6.0
...
4896	118.0	0.99540	3.07	0.50	9.4	6.0
4897	100.0	0.98938	3.06	0.44	11.8	6.0
4898	121.0	0.99074	3.24	0.46	10.6	6.0
4899	112.0	0.99298	3.29	0.54	9.7	5.0
4900	111.0	0.99254	2.99	0.46	9.4	6.0

[4901 rows x 12 columns]

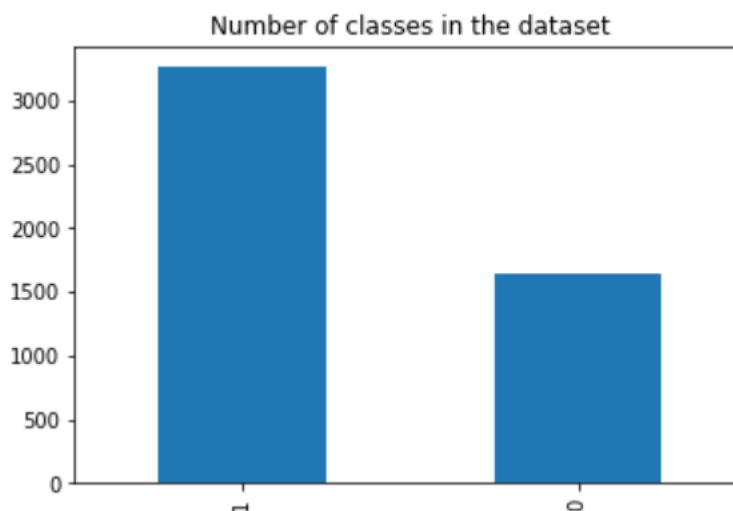
➔ `data['new_quality'] = [0 if i <= 5 else 1 for i in data.quality]`
 ➔ `print(data1)` *# Displaying the data after adding new column 'new quality'.*

	fixedacid	volacid	citricacid	residualsugar	chlorides	freesulfur	\
0	7.0	0.270	0.36	20.70	0.045	45.0	
1	7.2	0.230	0.32	8.50	0.058	47.0	
2	7.2	0.230	0.32	8.50	0.058	47.0	
3	7.0	0.270	0.36	20.70	0.045	45.0	
4	6.3	0.300	0.34	1.60	0.049	14.0	
...	
4896	4.9	0.235	0.27	11.75	0.030	34.0	
4897	6.1	0.340	0.29	2.20	0.036	25.0	
4898	5.7	0.210	0.32	0.90	0.038	38.0	
4899	6.5	0.230	0.38	1.30	0.032	29.0	
4900	6.5	0.240	0.19	1.20	0.041	30.0	

	totalsulfur	density	pH	sulphates	alcohol	quality	new_quality
0	170.0	1.00100	3.00	0.45	8.8	6.0	1
1	186.0	0.99560	3.19	0.40	9.9	6.0	1
2	186.0	0.99560	3.19	0.40	9.9	6.0	1
3	170.0	1.00100	3.00	0.45	8.8	6.0	1
4	132.0	0.99400	3.30	0.49	9.5	6.0	1
...
4896	118.0	0.99540	3.07	0.50	9.4	6.0	1
4897	100.0	0.98938	3.06	0.44	11.8	6.0	1
4898	121.0	0.99074	3.24	0.46	10.6	6.0	1
4899	112.0	0.99298	3.29	0.54	9.7	5.0	0
4900	111.0	0.99254	2.99	0.46	9.4	6.0	1

[4901 rows x 13 columns]

➔ `data['new_quality'].value_counts().plot(kind = 'bar', title = 'Number of classes in the dataset')` *# Plotting the graph for the number of classes found in the dataset*



```
➔ data1 = data.drop(columns = ['quality'])
➔ print(data1)    # Displaying the data after dropping the column 'quality'.
```

	fixedacid	volacid	citricacid	residualsugar	chlorides	freesulfur	\
0	7.0	0.270	0.36	20.70	0.045	45.0	
1	7.2	0.230	0.32	8.50	0.058	47.0	
2	7.2	0.230	0.32	8.50	0.058	47.0	
3	7.0	0.270	0.36	20.70	0.045	45.0	
4	6.3	0.300	0.34	1.60	0.049	14.0	
...	
4896	4.9	0.235	0.27	11.75	0.030	34.0	
4897	6.1	0.340	0.29	2.20	0.036	25.0	
4898	5.7	0.210	0.32	0.90	0.038	38.0	
4899	6.5	0.230	0.38	1.30	0.032	29.0	
4900	6.5	0.240	0.19	1.20	0.041	30.0	

	totalsulfur	density	pH	sulphates	alcohol	new_quality
0	170.0	1.00100	3.00	0.45	8.8	1
1	186.0	0.99560	3.19	0.40	9.9	1
2	186.0	0.99560	3.19	0.40	9.9	1
3	170.0	1.00100	3.00	0.45	8.8	1
4	132.0	0.99400	3.30	0.49	9.5	1
...
4896	118.0	0.99540	3.07	0.50	9.4	1
4897	100.0	0.98938	3.06	0.44	11.8	1
4898	121.0	0.99074	3.24	0.46	10.6	1
4899	112.0	0.99298	3.29	0.54	9.7	0
4900	111.0	0.99254	2.99	0.46	9.4	1

[4901 rows x 12 columns]

```
➔ data2 = data['new_quality']
➔ print(data2)    # Displaying the class values.
```

0	1
1	1
2	1
3	1
4	1
..	
4896	1
4897	1
4898	1
4899	0
4900	1

Name: new_quality, Length: 4901, dtype: int64

```
➔ print(pd.DataFrame(data2.value_counts())) # no of high classes and low classes.
```

	new_quality
1	3260
0	1641

```
➔ X_train, X_test, y_train, y_test = train_test_split(data1, data2, test_size=0.2)
➔ print(X_train)           # Displaying trained split data.
```

	fixedacid	volacid	citricacid	residualsugar	chlorides	freesulfur	\
745	6.8	0.21	0.36	18.10	0.046	32.0	
4394	6.7	0.21	0.34	1.50	0.035	45.0	
2507	7.1	0.85	0.49	8.70	0.028	40.0	
4107	6.3	0.24	0.29	13.70	0.035	53.0	
2567	7.6	0.18	0.49	18.05	0.046	36.0	
...	
3721	7.1	0.14	0.33	1.00	0.104	20.0	
3027	7.7	0.39	0.34	10.00	0.056	35.0	
2524	6.6	0.30	0.24	1.20	0.034	17.0	
771	6.9	0.13	0.28	13.30	0.050	47.0	
548	7.7	0.44	0.24	11.20	0.031	41.0	

	totalsulfur	density	pH	sulphates	alcohol	new_quality
745	133.0	1.00000	3.27	0.48	8.8	0
4394	123.0	0.98949	3.24	0.36	12.6	1
2507	184.0	0.99620	3.22	0.36	10.7	0
4107	134.0	0.99567	3.17	0.38	10.6	1
2567	158.0	0.99960	3.06	0.41	9.2	0
...
3721	54.0	0.99057	3.19	0.64	11.5	1
3027	178.0	0.99740	3.26	0.60	10.2	0
2524	121.0	0.99330	3.13	0.36	9.2	0
771	132.0	0.99655	3.34	0.42	10.1	1
548	167.0	0.99480	3.12	0.43	11.3	1

[3920 rows x 12 columns]

```
➔ accuracy = {}
➔ knn = KNeighborsClassifier(n_neighbors = k)
➔ model = knn.fit(X_train,y_train)
➔ pred = knn.predict(X_test)
➔ accuracy = metrics.accuracy_score(y_test,pred)
➔ print(accuracy)           # calculating the accuracy.
```

0.7145769622833843

```
➔ kfold = KFold(n_splits=5, shuffle=False)
➔ cv_score = cross_val_score(knn, data1, data2,cv = cv1), scoring='accuracy')
➔ print(cv_score)           # calculating accuracy for each fold.
```

[0.67686035 0.67142857 0.65816327 0.7122449 0.69489796]

```
➔ print(cv_score.mean())    # displaying mean accuracy.
```

0.6827190080925336

- ➔ `grid = KNeighborsClassifier()`
- ➔ `grid1 = {'n_neighbors': np.arange(R1, R2)}`
- ➔ `knn_gridcv = GridSearchCV(grid, grid1, cv = cv1)`
- ➔ `knn_gridcv.fit(data1, data2)` *# calculating accuracy for each value of k in the range 1-25.*

```
GridSearchCV(cv=5, error_score='raise-deprecating',
             estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                           metric='minkowski',
                                           metric_params=None, n_jobs=None,
                                           n_neighbors=5, p=2,
                                           weights='uniform'),
             iid='warn', n_jobs=None,
             param_grid={'n_neighbors': array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,
18, 19, 20, 21, 22, 23, 24])},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)
```

- ➔ `print(knn_gridcv.best_params_)` *# the best value of k.*

```
{'n_neighbors': 1}
```

- ➔ `print(knn_gridcv.best_score_)` *# the accuracy for gridsearchCV*

```
0.7463782901448683
```